

Dating with constraints

A detailed tutorial on McmcDate

Dominik Schrempf

January 13, 2023

Contents

1	Definitions	2
2	Dating a tree	2
2.1	Node age calibrations	3
2.2	Node order constraints	3
2.3	Node braces	4
3	Pipeline	4
4	Example analysis	5
4.1	Step 1: Alignment and unrooted topology	5
4.2	Step 2: Phylogenetic inference	6
4.3	Step 3: Rooted topology and auxiliary data	6
4.4	Step 4: Dating with McmcDate	6
A	Auxiliary data specifications	7
A.1	Node age calibrations	7
A.2	Node order constraints	7
A.3	Node braces	7
B	Internals	7
B.1	Haskell modules	8
B.1.1	More important modules	8
B.1.2	Other modules	8

Abstract

McmcDate can date a phylogenetic tree with constraints. That is, it can estimate the ages of the ancestral nodes of a phylogenetic tree with node age calibrations, node order constraints, and node braces. McmcDate is fast because it approximates the phylogenetic likelihood with a multivariate normal distribution.

Here, I briefly review the process of dating phylogenetic trees and provide an example dating analysis of a tree of eukaryotes. I use data from (Strassert et al. 2021) who provide the alignment and node age calibrations from fossils. Moreover, I use node order constraints from two possible successions of plastid endosymbiosis events.

This tutorial is a work in progress.

1 Definitions

We use the term *(rooted) tree* to denote a directed acyclic graph with node labels and branch lengths, and in which not more than one branch connects any two nodes. Usually, nodes correspond to (ancestral) species, and branch lengths to sequence distance or time. We use the term *(rooted) topology* to denote a (rooted) tree without information about branch lengths. We put the word *rooted* into parentheses, because it is sometimes omitted. The root of the tree or topology is the only node with all branches pointing away to other nodes, and is the oldest node. Leaves are extant nodes with a single branch pointing towards them.

For a given topology, *dating the tree*¹ corresponds to finding branch lengths for this topology which are measured in absolute time units (for example, in Million years), and which describe the data in the *best* way. *Best* can mean different things, and we will carefully analyze what *best* means in our case. Sometimes, all we can do is date a tree with relative time units (Section 2.1).

Similarly, we use the term *unrooted tree* to denote an undirected acyclic graph with node labels and branch lengths, and in which not more than one branch connects any two nodes. Finally, we use the term *unrooted topology* to denote an unrooted tree without information about branch lengths. We need unrooted trees and topologies because we will use reversible substitution models (for a review, see Yang 2006) to infer unrooted trees with branch lengths measured in expected number of substitutions. Reversible substitution models are unable to discriminate between the two directions in time.

2 Dating a tree

Dating a tree is difficult because we want to estimate branch lengths measured in absolute time units. However, an alignment only contains information about distances measured in expected number of substitutions. In particular, it does not contain information about the evolutionary rates nor about the actual elapsed absolute time units. That is, for a given branch,

$$\begin{aligned} d \text{ in [expected number of substitutions]} = \\ r \text{ in [expected number of substitutions per year]} \cdot t \text{ in [years]}, \end{aligned} \quad (1)$$

where d is the branch length measured in expected number of substitutions, r is the evolutionary rate on this branch, and t is the branch length measured in absolute time units. The situation is severe because r and t are confounded! If we multiply t with a positive number c and divide r by the same positive number c , the distance d stays constant. Since the likelihood only depends on the distance d we can not discriminate between parameters

$$t' = t \cdot c, \quad (2)$$

$$r' = r/c, \quad (3)$$

¹We should probably say *dating the topology* but this phrase is not used.

for *any* $c \in (0, \infty)$! This means we need auxiliary data to constrain the ages of at least some internal nodes. There are three forms of constraints from auxiliary data which are explained in the next sections: (a) node age calibrations, (b) node order constraints, and (c) node braces. Appendix A details the specifications about how auxiliary data is used with McmcDate.

2.1 Node age calibrations

We call constraints on the ages of internal nodes *node age calibrations*. Without node age calibrations, all dated trees are equally likely! Since most phylogenetic dating methods are Bayesian, and impose prior functions on all of their parameters, the values of the posterior function of different trees still differ in a deceiving way. However, in the absence of node age calibrations, the differences in the values of the posterior function are exclusively caused by the chosen prior functions, and not by the information we have obtained from the data.

Even more, in order to limit the maximum ages of the nodes, we require at least one node age calibration with a maximum boundary. The closer in terms of absolute time units the calibrated node with a maximum boundary is to the root, the better. Ideally, we have a maximum age boundary for the root itself, but sometimes the root age boundaries are unknown. Section 4.3 provides details on how to add node age calibrations to your dating analysis with McmcDate.

I can not repeat this often enough: **Without proper node age calibrations, and in particular, without at least one maximum age boundary, dating a tree in absolute time units is impossible.** Without a single maximum node age boundary, we can still date the tree using branch lengths measured in *relative time units*. In particular, we achieve this by fixing the tree height to be 1.0. McmcDate automatically falls back to dating in relative time when no maximum node age boundary is found. If other programs date trees in absolute time units without a maximum age boundary, do not trust the results!

2.2 Node order constraints

Next to calibrating the absolute ages of nodes, we can also constrain the relative order of nodes on a tree. For example, due to external analyses, we may have detailed knowledge about a horizontal gene transfer. That is, we do know the exact donor and the recipient branches. In this case, a direct horizontal gene transfer can only happen, if these branches coexist for a period of time (Figure 1). If the branches do not coexist for a period of time but the recipient branch is younger than the donor branch, the horizontal gene transfer can still happen indirectly through multiple events. An indirect gene transfer requires an intermediate lineage carrying the gene. The intermediate lineage has either gone extinct (Szöllősi et al. 2013), has not been sampled, or has lost the gene. Gene transfers are impossible if the recipient branch is younger than the donor branch.

Let the ages of the old and young nodes of the donor and recipient branches be DO , DY , RO , and RY , respectively. Then, a direct horizontal gene transfer provides us with two node order constraints,

$$DY < RO, \text{ and } DO > RY, \quad (4)$$

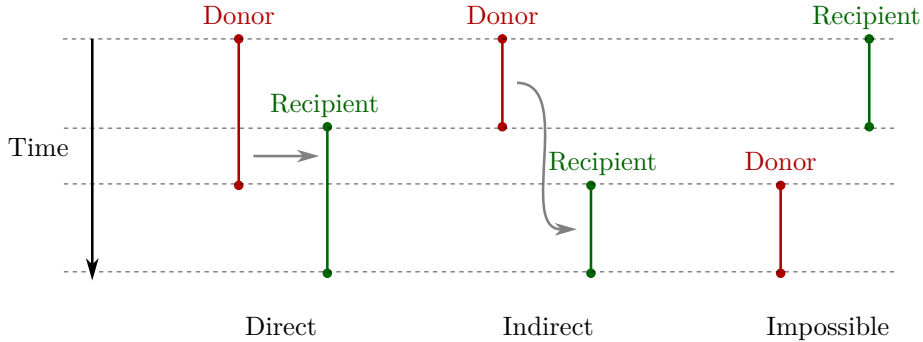


Figure 1: Direct, indirect and impossible horizontal gene transfers. A direct horizontal gene transfer happens between branches coexisting in time. An indirect horizontal gene transfer involves multiple events and an intermediate lineage. If the recipient branch is older than the donor branch, a horizontal gene transfer is impossible.

where $<$ means *younger than* and $>$ means *older than*. If we allow indirect gene transfers, we only get one node order constraint,

$$DO > RY. \quad (5)$$

The last constraint is enough to prohibit impossible gene transfers as depicted in Figure 1.

Sometimes, we do not know the exact donor branch but only that the donor branch must be on a subtree. However, we usually do know the recipient branch which is the stem of the subtree comprising the species which contain the horizontally transferred gene. In this case, the older node of the stem of the donor subtree has to be older than the younger node of the recipient branch.

2.3 Node braces

Finally, sometimes we know that two or more nodes have the same age. For example, when analyzing gene trees, an early gene duplication event may separate the gene tree into two subtrees. Subsequent speciation events may be observable on both subtrees in that the corresponding speciation nodes should have similar ages. In this tutorial we will not brace nodes.

3 Pipeline

McmcDate is fast because it uses a clever way to approximate the phylogenetic likelihood similar to McmcTree (Yang and Rannala 2005). McmcTree approximates the surface of the phylogenetic likelihood using Taylor expansion (Reis and Yang 2011). That is, in a prior analysis we compute the maximum likelihood together with the gradient and the Hessian matrix which we use in a subsequent analysis to date the tree.

Here, we use a similar technique, albeit with a Bayesian method to estimate the posterior distribution of branch lengths measured in expected number substitutions. In particular, the pipeline for dating a tree with McmcDate is:

1. Prepare a multi sequence alignment and an unrooted topology.

2. For this alignment and unrooted topology, infer a distribution of unrooted trees with branch lengths measured in expected number of substitutions.
3. Find a rooted topology to date. Prepare auxiliary data such as node age calibrations or node order constraints.
4. Date the topology with MmcDate. In particular, infer a distribution of rooted trees with branch lengths measured in absolute time units. If the node age calibrations are insufficient, we can still use relative time units.

Steps 3 and 4 are fast, especially when compared to Step 2. Hence, we can conveniently amend parameters used in the actual dating analysis, or compute dated trees for different roots. The following section shows an example analysis and describes the steps of the pipeline in detail.

4 Example analysis

4.1 Step 1: Alignment and unrooted topology

We are going to date a topology of eukaryotes (Strassert et al. 2021). The full published data set contains three alignments:

1. The alignment with the highest number of taxa comprises 733 eukaryotes with 62723 amino acids.
2. The authors used the first alignment to infer a tree with IQ-TREE (Minh et al. 2020) so they can filter taxa in an informed way (for example, keep slow evolving taxa). The reduced data set comprises 136 operational taxonomic units with 73460 amino acids.
3. A small data set with 63 operational taxonomic units and 73460 amino acids for tree inference with Bayesian models.

Here, we focus on the second alignment with 136 taxa. We use shorter taxon names and I have converted the file from FASTA format to Phylip format which is required by Phylobayes (Lartillot et al. 2013). I provide the alignment `strassert-136taxa.phy` in the [data subfolder](#) of this tutorial:

```
1 data/strassert-136taxa.phy
```

For the phylogenetic inference in the next step, we also need an unrooted topology. I have unrooted the tree in Figure 3 in Strassert et al. (2021). Similar to the alignment, I provide the unrooted topology in the [data subfolder](#) of this tutorial:

```
1 data/strassert-136taxa.unrooted.tre
```

Please also see the [README](#) in the [data subfolder](#).

4.2 Step 2: Phylogenetic inference

Here, we use a Bayesian method to estimate the posterior distribution. In particular,

- Use Phylobayes (Lartillot et al. 2013).
- Decide on evolutionary model depending on the size of the data set and the computational requirements. Recommended models from preferred but slow and complex to fast and simple: GTR+CAT+G4, LG+CAT+G4, LG+EDM64+G4, LG+C60+G4, LG+G4.

We specify an evolutionary model with exchangeabilities EX, and across-site compositional heterogeneity model ASCH as EX+ASCH. All discussed evolutionary models used for simulations as well as inferences implicitly use discrete gamma rate heterogeneity with four components.

- GTR model (Tavaré 1986).
- CAT model (Lartillot and Philippe 2004).
- Gamma rate variation model (Yang 1993).
- LG model (Le and Gascuel 2008).
- EDM model (Schrempf et al. 2020).
- C60 model (Quang et al. 2008)

4.3 Step 3: Rooted topology and auxiliary data

Strassert et al. (2021) discuss one unrooted topology with two possible root positions. (1) The root separates amorphea from diaphoretickes and excavates (Figure 3 in Strassert et al. 2021), and (2) the root separates amorphea and excavates from diaphoretickes. Here, we choose option (1) which is the more plausible one.

In this tutorial, we use 33 fossil calibrations compiled by Strassert et al.

- Node order calibrations (Yang and Rannala 2005).
- Relative node order constraints (Szöllősi et al. 2022).
- McmcDate can also brace nodes (Appendix A).

Hello.

4.4 Step 4: Dating with McmcDate

- McmcDate is a Haskell program; Appendix B provides details about the internals of McmcDate.

Hello.

A Auxiliary data specifications

The specifications match McmcDate version 1.0.0.0 and may change between different versions of McmcDate.

A.1 Node age calibrations

Node age calibrations can be provided in two ways:

- with a comma separated values (CSV) file, or
- with a Newick tree file (MCMCTree specification; see the documentation of MCMCTree; only L, U, and B are supported).

If the file ends with `csv`, assume the calibrations are provided in CSV format. Otherwise, assume the calibrations are provided on a Newick tree. The CSV file has rows of the following format:

```
1 CalibName,LeafA,LeafB,LowerBoundary,LowerBoundaryProbMass,UpperBoundary,UpperBoundaryProbMass
```

In this case, the calibrated node is uniquely defined as the most recent common ancestor of `LeafA` and `LeafB`. The age of the node is calibrated between the lower and upper boundary. The probability mass describes the softness (or hardness) of a boundary. In other words, the probability mass describes the steepness of the decline of the prior function outside the calibration interval. In general, the larger the probability mass the softer the boundary. We specify the probability mass with respect to a normalized time interval of size 1.0. That is, probability masses have to be strictly positive and less than 1.0, which is the total probability mass in the unit interval.

I usually use values between 0.0001 (hard) and 0.03 (soft). If unsure, use probability masses of 0.025, which corresponds to 2.5 percent probability at each boundary or constraint. A probability mass close to 1.0 will correspond to a prior function too soft to have any effect. Note that this way of specifying boundary softness using relative values independent of the actual node ages differs from MCMCTree which uses absolute values (Yang and Rannala 2005). When using a Newick tree to specify node age calibrations, and no probability masses are provided, a default value of 0.01 is used. This measure is in place to support the same input files as MCMCTree does.

To specify one-sided node age calibrations, omit the other boundary and the corresponding probability mass. For example, the following line defines a node age calibration with a lower boundary at `1e6` time units (years in this case) with probability mass 0.025:

```
1 Primates,Human,Chimpanzees,1e6,0.025,,
```

A.2 Node order constraints

A.3 Node braces

B Internals

McmcDate is a Haskell program.

Recommend cabal, but there is also stack (-s) option.

The `wrapper script` used in this tutorial tries to make a good compromise between usability and customizability. It exposes some, but not all functionality of McmcDate. Most notably,

- Based on `mcmc`.
- Based on `elynx-tree`.
- Explain code a bit (I guess mostly proposals).

B.1 Haskell modules

Modules containing definitions specific to the analysis are in the `app` subfolder of the McmcDate repository.

B.1.1 More important modules

Definitions Proposals and monitors, configuration.

State State space. If you try to understand what is going on, or if you want to change analysis settings, this should be your starting point.

B.1.2 Other modules

Hamiltonian Hamiltonian proposal.

Main Functions to prepare the data, run and continue the Metropolis-Hastings-Green algorithm.

Monitor Prior specific monitoring functions.

Options Handle command line options.

Probability Prior and likelihood functions.

Tools Miscellaneous tools.

References

- Lartillot, N. and H. Philippe (2004). “A Bayesian mixture model for across-site heterogeneities in the amino-acid replacement process.” In: *Molecular Biology and Evolution* 21.6, pp. 1095–1109. DOI: [10.1093/molbev/msh112](https://doi.org/10.1093/molbev/msh112).
- Lartillot, N., N. Rodrigue, D. Stubbs, and J. Richer (2013). “PhyloBayes MPI: Phylogenetic Reconstruction with Infinite Mixtures of Profiles in a Parallel Environment.” In: *Systematic Biology* 62.4, pp. 611–615. DOI: [10.1093/sysbio/syt022](https://doi.org/10.1093/sysbio/syt022).
- Le, S. Q. and O. Gascuel (2008). “An improved general amino acid replacement matrix.” In: *Molecular Biology and Evolution* 25.7, pp. 1307–1320. DOI: [10.1093/molbev/msn067](https://doi.org/10.1093/molbev/msn067).

- Minh, B. Q., H. A. Schmidt, O. Chernomor, D. Schrempf, M. D. Woodhams, A. von Haeseler, and R. Lanfear (2020). “IQ-TREE 2: New models and efficient methods for phylogenetic inference in the genomic era.” In: *Molecular Biology and Evolution* 37.5, pp. 1530–1534. DOI: [10.1093/molbev/msaa015](https://doi.org/10.1093/molbev/msaa015).
- Quang, L. S., O. Gascuel, and N. Lartillot (2008). “Empirical profile mixture models for phylogenetic reconstruction.” In: *Bioinformatics* 24.20, pp. 2317–2323. DOI: [10.1093/bioinformatics/btn445](https://doi.org/10.1093/bioinformatics/btn445).
- Reis, M. dos and Z. Yang (2011). “Approximate Likelihood Calculation on a Phylogeny for Bayesian Estimation of Divergence Times.” In: *Molecular Biology and Evolution* 28.7, pp. 2161–2172. DOI: [10.1093/molbev/msr045](https://doi.org/10.1093/molbev/msr045).
- Schrempf, D., N. Lartillot, and G. Szöllősi (2020). “Scalable empirical mixture models that account for across-site compositional heterogeneity.” In: *Molecular Biology and Evolution*. DOI: [10.1093/molbev/msaa145](https://doi.org/10.1093/molbev/msaa145).
- Strassert, J. F. H., I. Irisarri, T. A. Williams, and F. Burki (2021). “A molecular timescale for eukaryote evolution with implications for the origin of red algal-derived plastids.” In: *Nature Communications* 12.1. DOI: [10.1038/s41467-021-22044-z](https://doi.org/10.1038/s41467-021-22044-z).
- Szöllősi, G. J., S. Höhna, T. A. Williams, D. Schrempf, V. Daubin, and B. Boussau (2022). “Relative Time Constraints Improve Molecular Dating.” In: *Systematic Biology* 71.4, pp. 797–809. DOI: [10.1093/sysbio/syab084](https://doi.org/10.1093/sysbio/syab084).
- Szöllősi, G. J., E. Tannier, N. Lartillot, and V. Daubin (2013). “Lateral Gene Transfer from the Dead.” In: *Systematic Biology* 62.3, pp. 386–397. DOI: [10.1093/sysbio/syt003](https://doi.org/10.1093/sysbio/syt003).
- Tavaré, S. (1986). “Some Probabilistic and Statistical Problems in the Analysis of DNA Sequences.” In: *Lectures on Mathematics in the Life Sciences* 17, pp. 57–86.
- Yang, Z. (1993). “Maximum-likelihood estimation of phylogeny from DNA sequences when substitution rates differ over sites.” In: *Molecular Biology and Evolution*. DOI: [10.1093/oxfordjournals.molbev.a040082](https://doi.org/10.1093/oxfordjournals.molbev.a040082).
- Yang, Z. and B. Rannala (2005). “Bayesian Estimation of Species Divergence Times Under a Molecular Clock Using Multiple Fossil Calibrations with Soft Bounds.” In: *Molecular Biology and Evolution* 23.1, pp. 212–226. DOI: [10.1093/molbev/msj024](https://doi.org/10.1093/molbev/msj024).
- Yang, Z. (2006). *Computational molecular evolution*. Vol. 284.